



SAPIENZA
UNIVERSITÀ DI ROMA

Self-Driving Car 3D Model with Pattern Recognition via a Convolutional Neural Network

COMPUTER VISION PRESENTATION BY MATTEO PRATA & ANTONIO RICCIARDI - 11/01/2019

Introduction to the problem

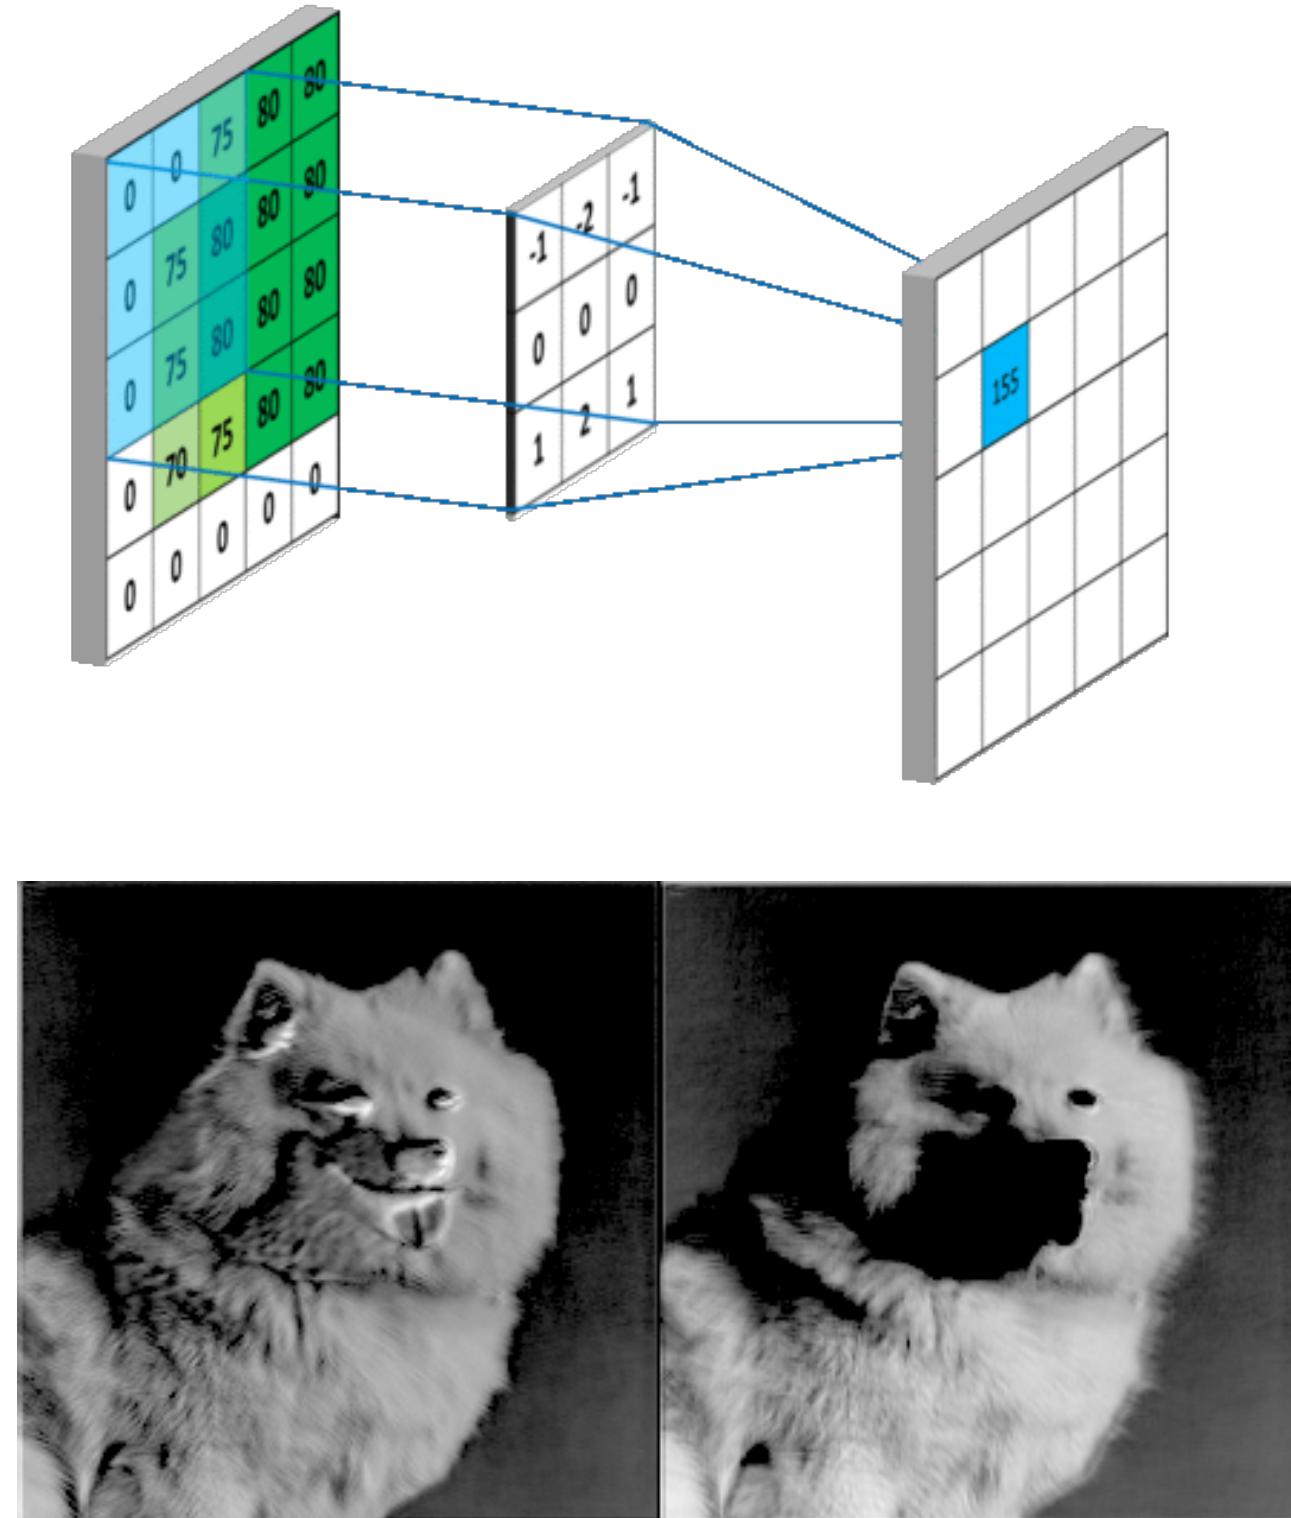
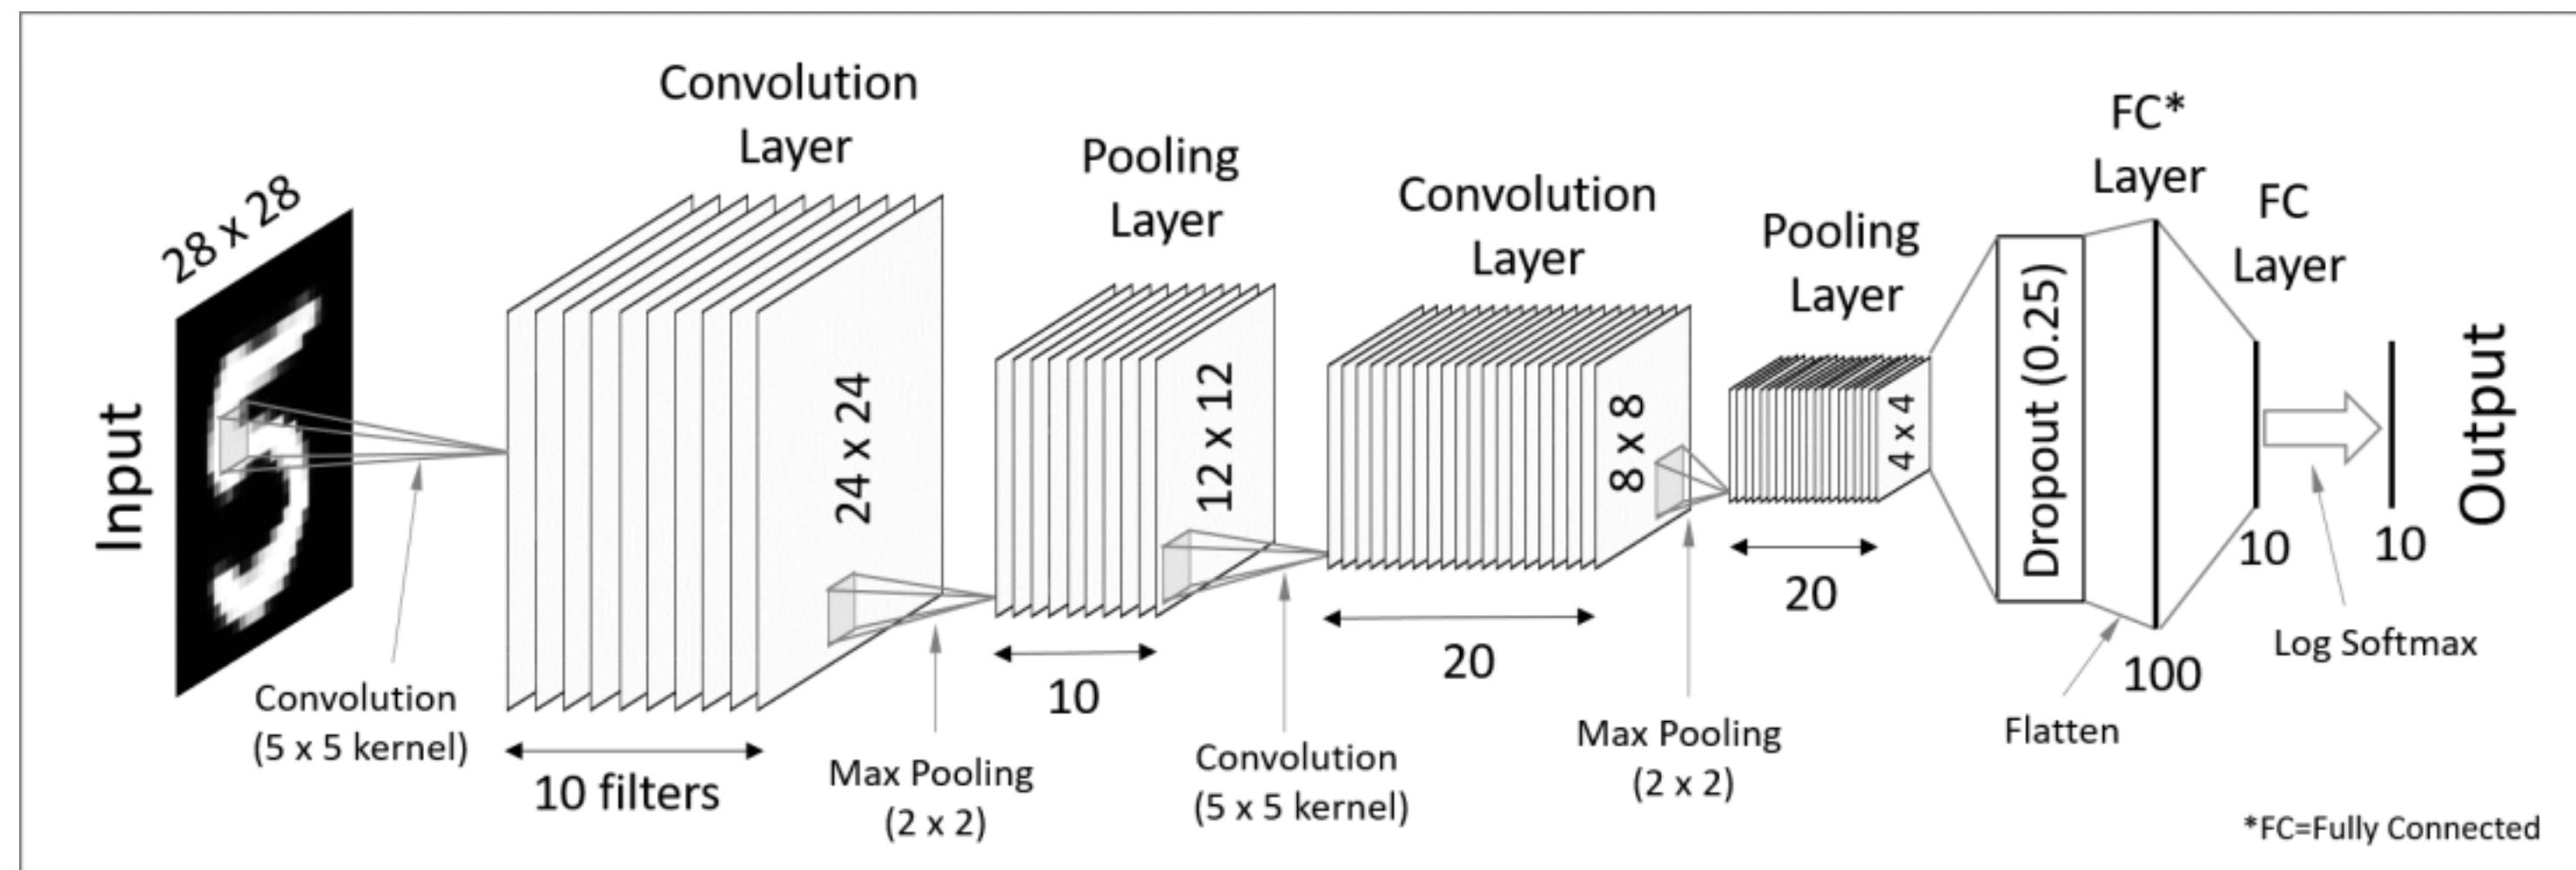
“End to End Learning for Self-Driving Cars”

By Marius Bojark et al. (2016)

- Train a Convolutional Neural Network (CNN) to map raw pixels, from a front-facing camera mounted on a car, directly to the steering commands
- CNN goes beyond pattern recognition, it learns the entire processing pipeline to steer the car
- Empirical demonstration that CNNs can learn complex tasks of road following, only through vision and without any centralized set of manual constraints

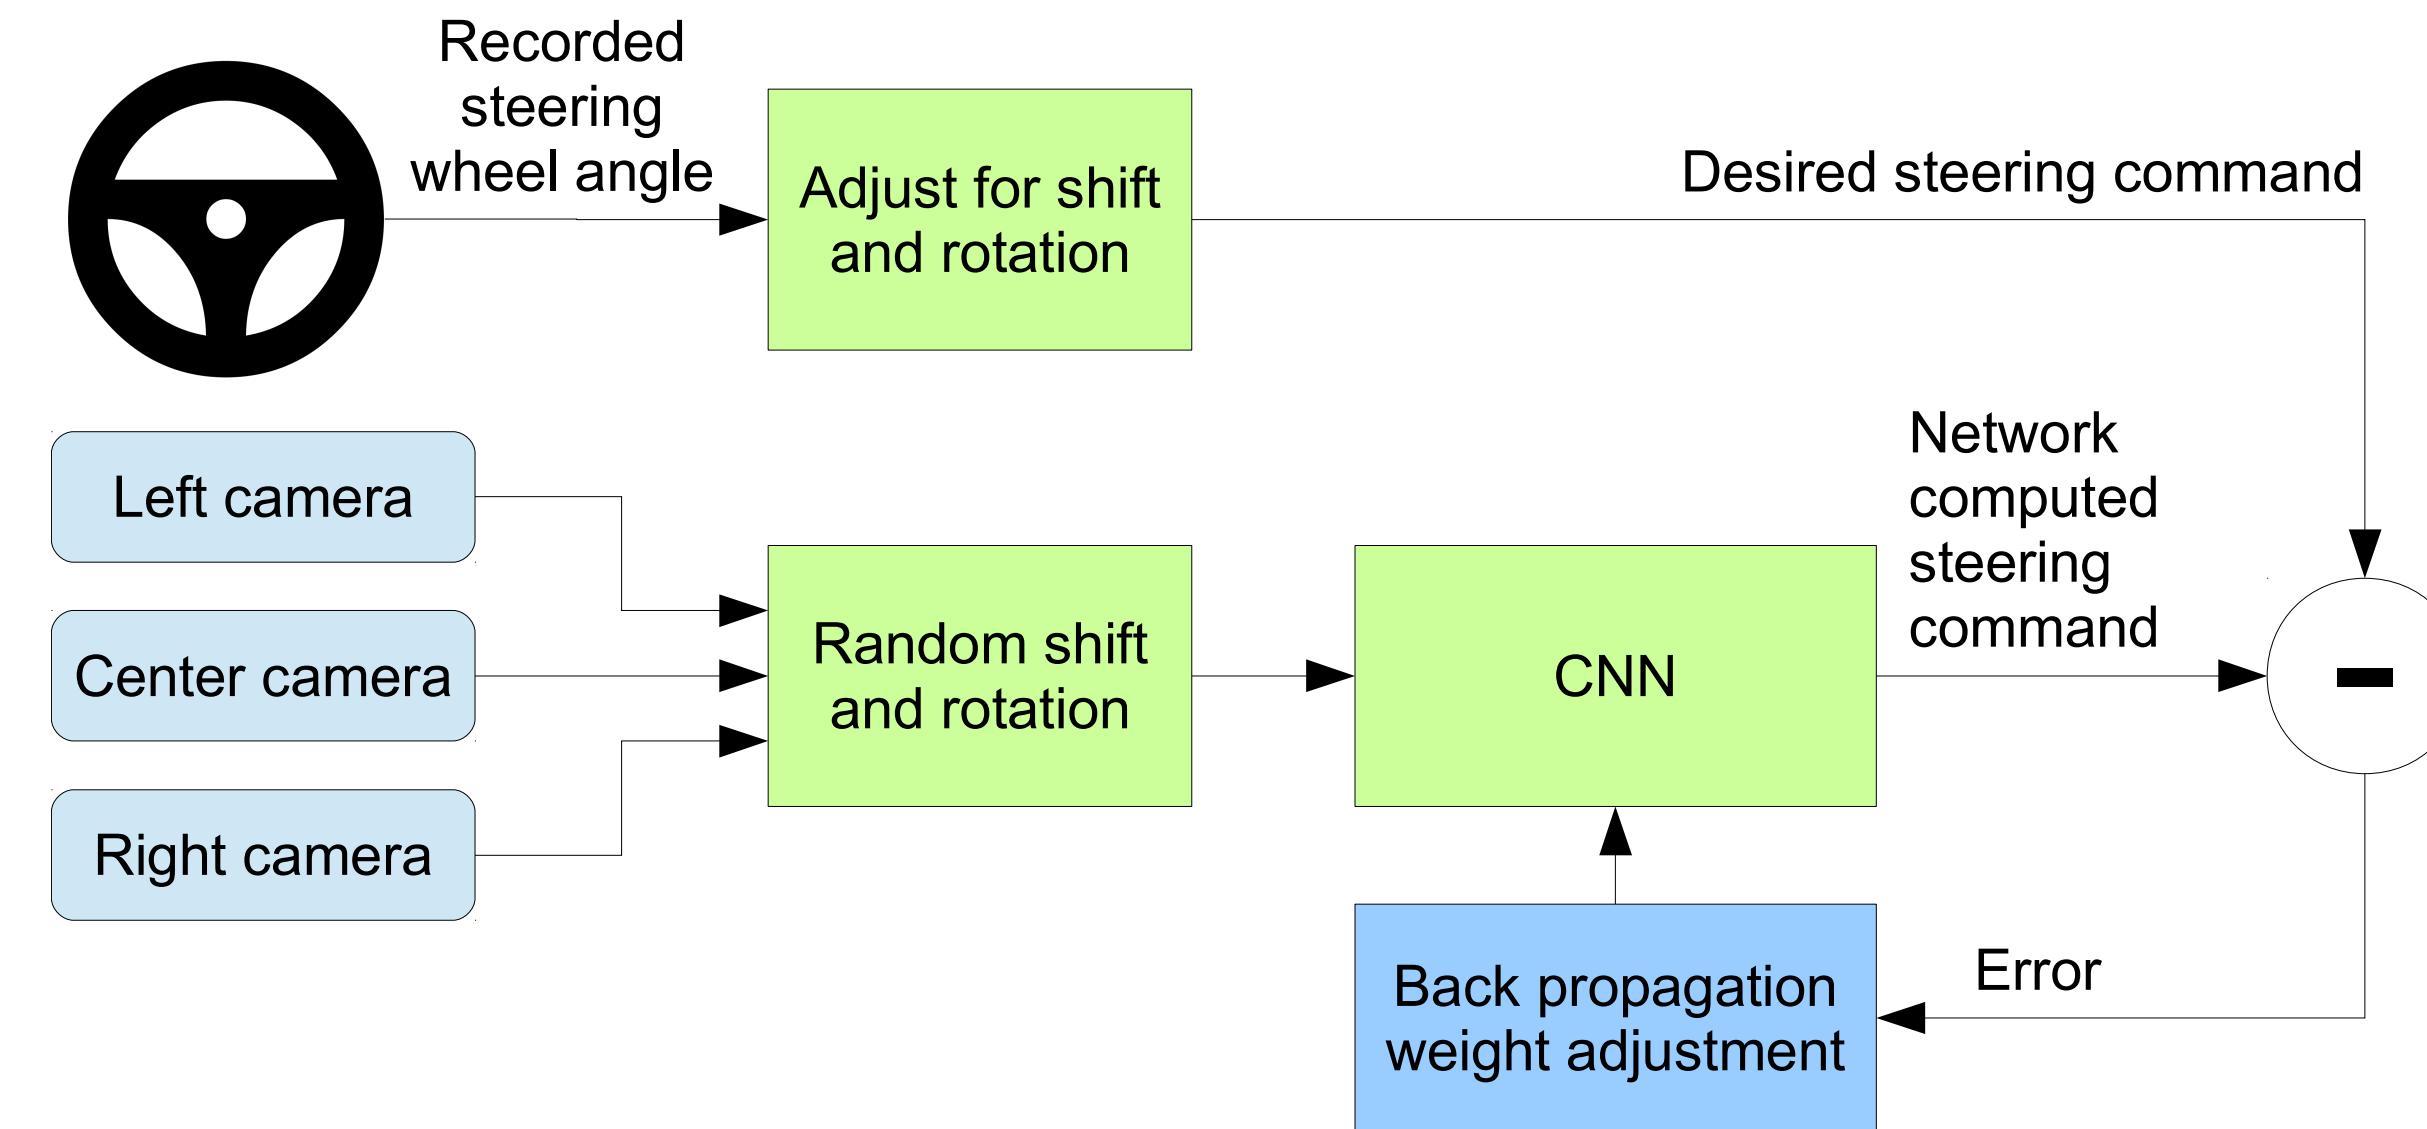
The CNNs power

- CNNs are feed forward neural networks which have revolutionized pattern recognition
- Features are learned automatically from training examples
- Convolution captures the 2D nature of images



End to End Learning for Self-Driving Car

- **PHASE 1**
Training phase

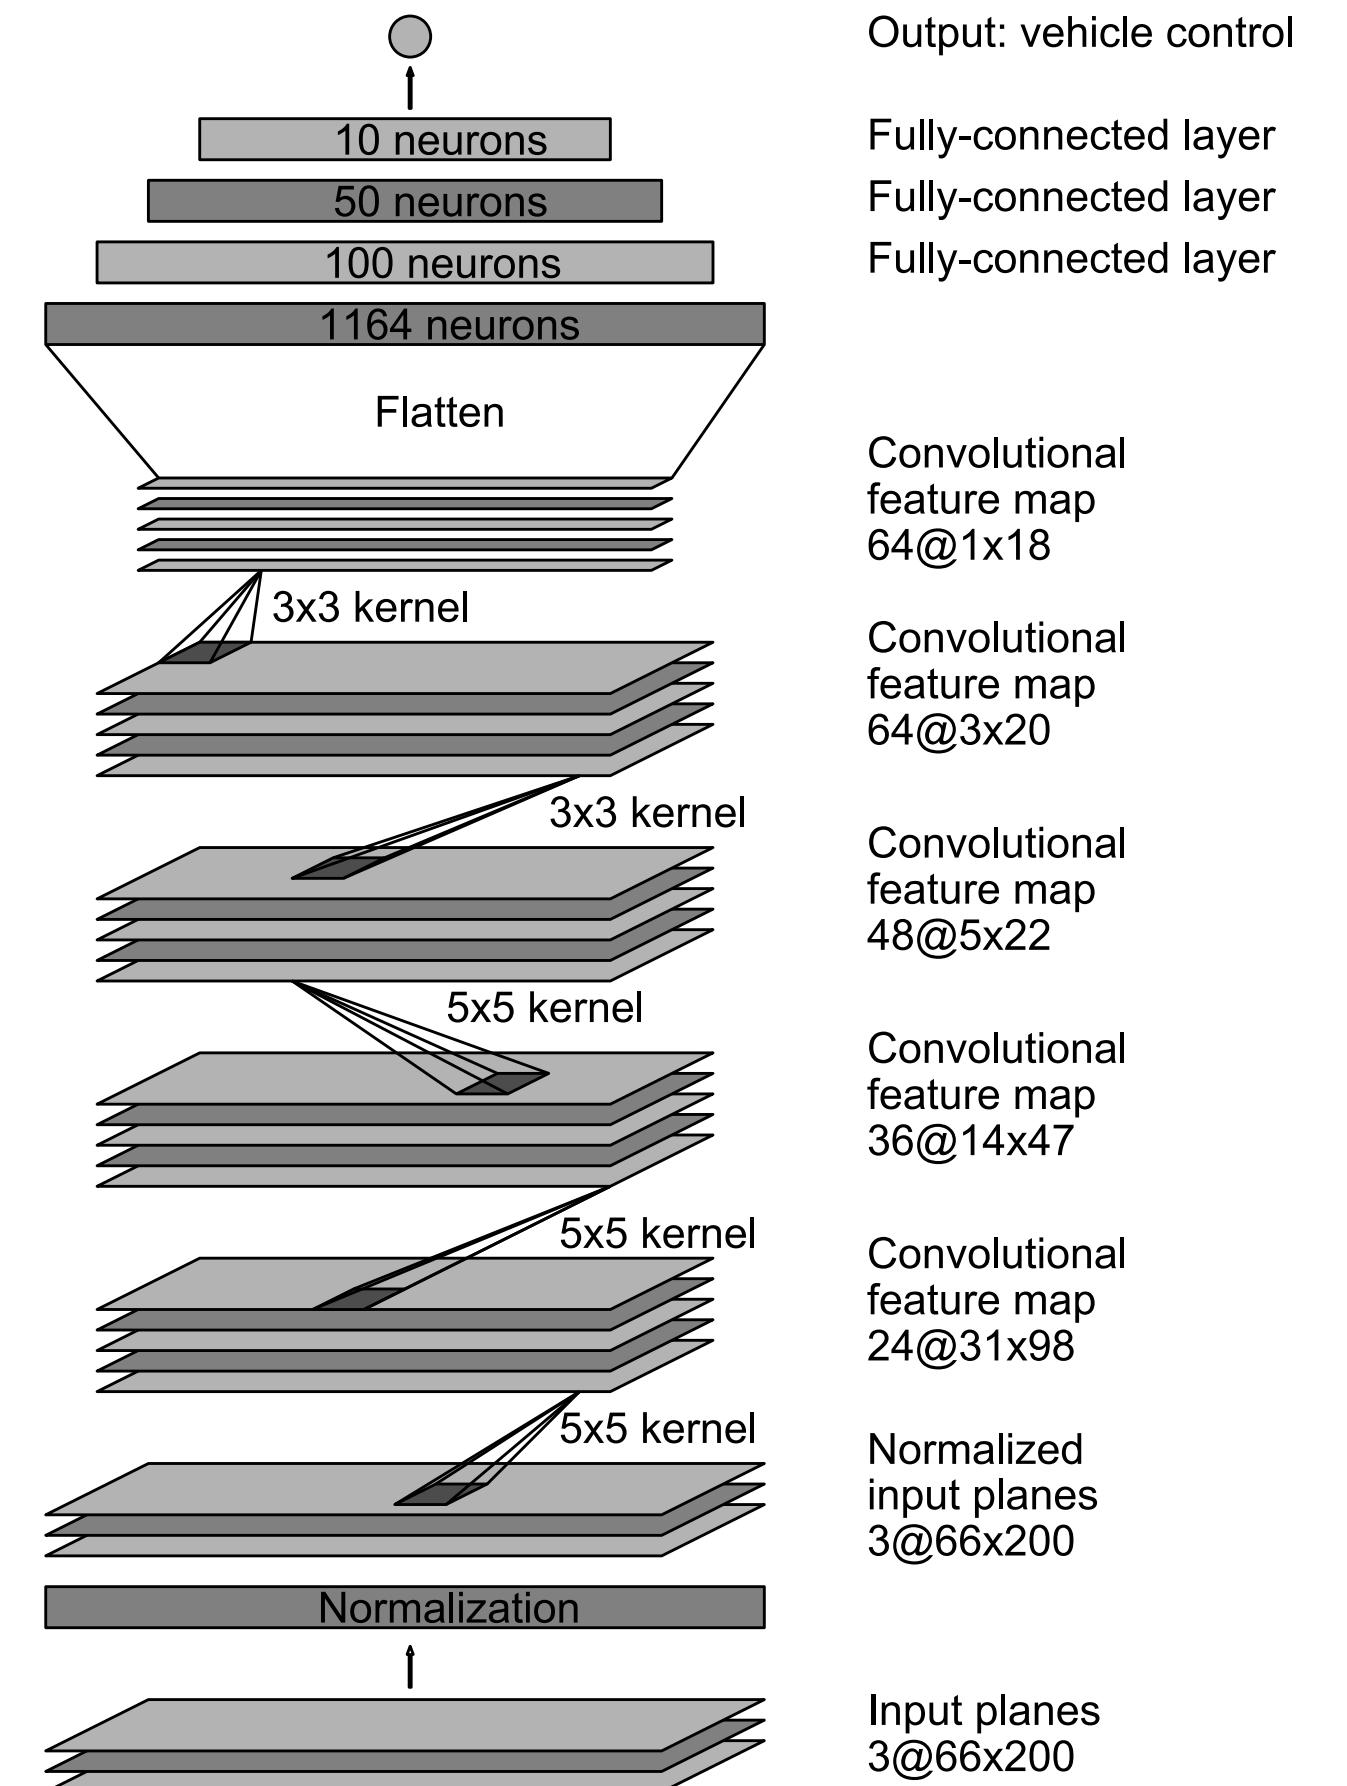


- **PHASE 2**
Test phase



CNN architecture

- Input images captured from a wide variety of roads, diverse lights and weather conditions
- Input images are split into YUV planes and passed through the CNN
- **9 layers:** 1 normalization, 5 convolutional, 3 fully connected layers.
- 27 million connections and 250 thousand parameters.
- Train the weights of the FC layer so to minimize the mean squared error of predicted steering command and the correct one



Data selection and augmentation

- Sampled **10 FPS** from the videos recorded from the three cameras, avoid useless similar frames
- Sample only the frames where the driver was staying in a lane, discard the rest
 - Added artificial **flips** and **shifts** to let the model learn how to recover from bad orientations, magnitude chosen from a normal distribution



Model evaluation

- The **autonomy** estimates what percentage of time the network could drive the car

$$\text{autonomy} = \left(1 - \frac{(\text{number of interventions}) \cdot 6 \text{ seconds}}{\text{elapsed time [seconds]}}\right) \cdot 100$$

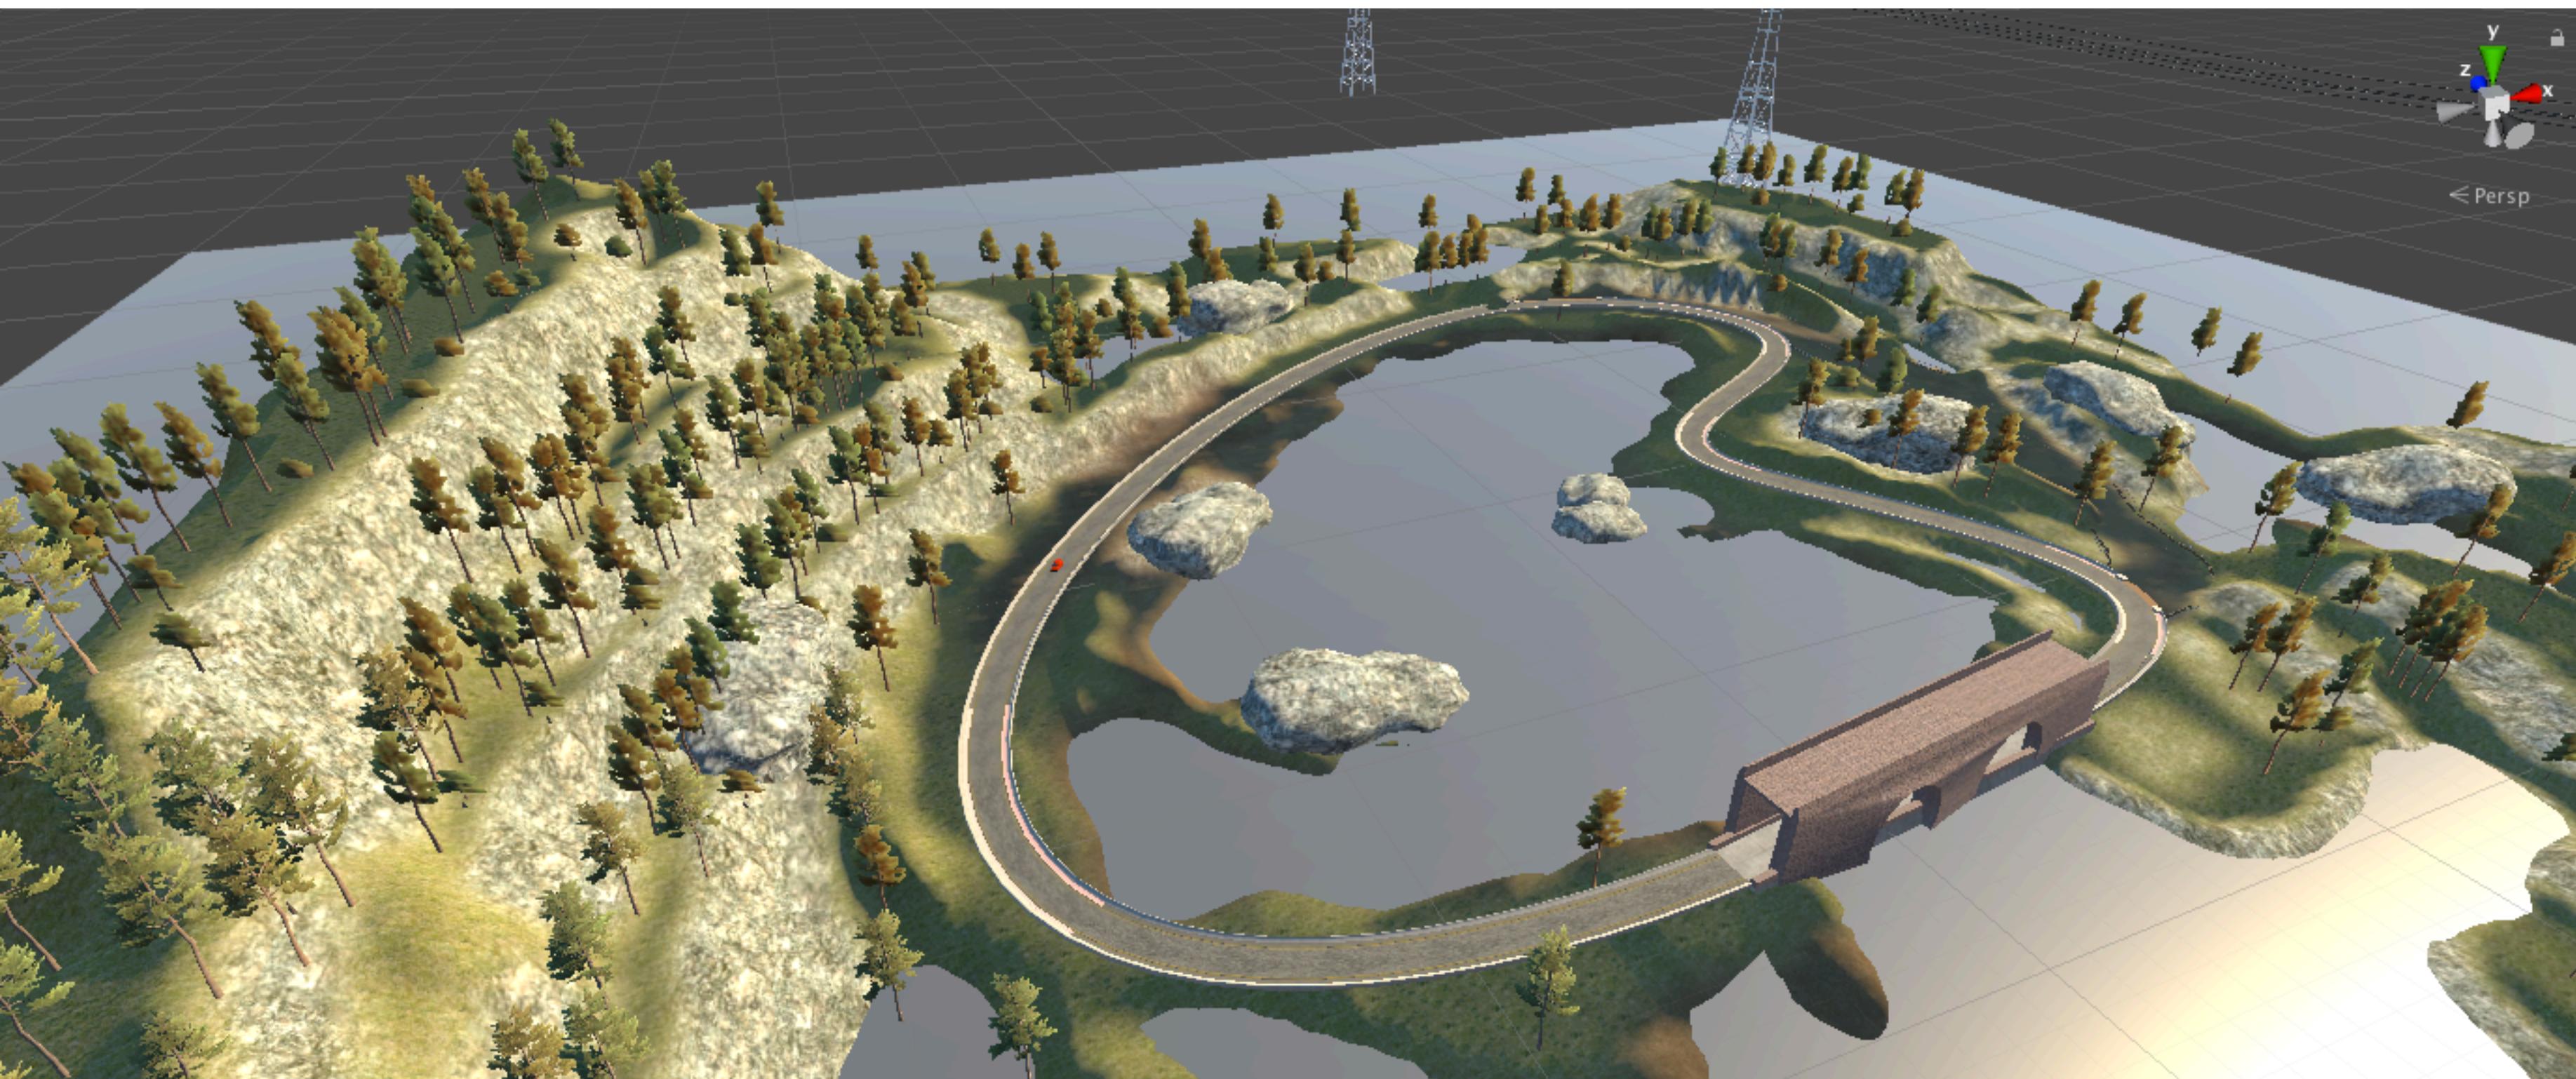
- Autonomy increases if the ratio between human interventions and elapsed time decreases
- If we had 10 interventions in 10 minutes, we would have 90% of autonomy
- For a typical drive in Monmouth Country NJ, the car was autonomous approximately **98%** of the time. On a multi-lane divided highway in Garden State Parkway, **100%**.

Demo on the road



Self-Driving 3D car driven by a CNN

- The 3D environment in Unity



credits: <https://assetstore.unity.com/packages/3d/environments/roadways/lake-race-track-55908>

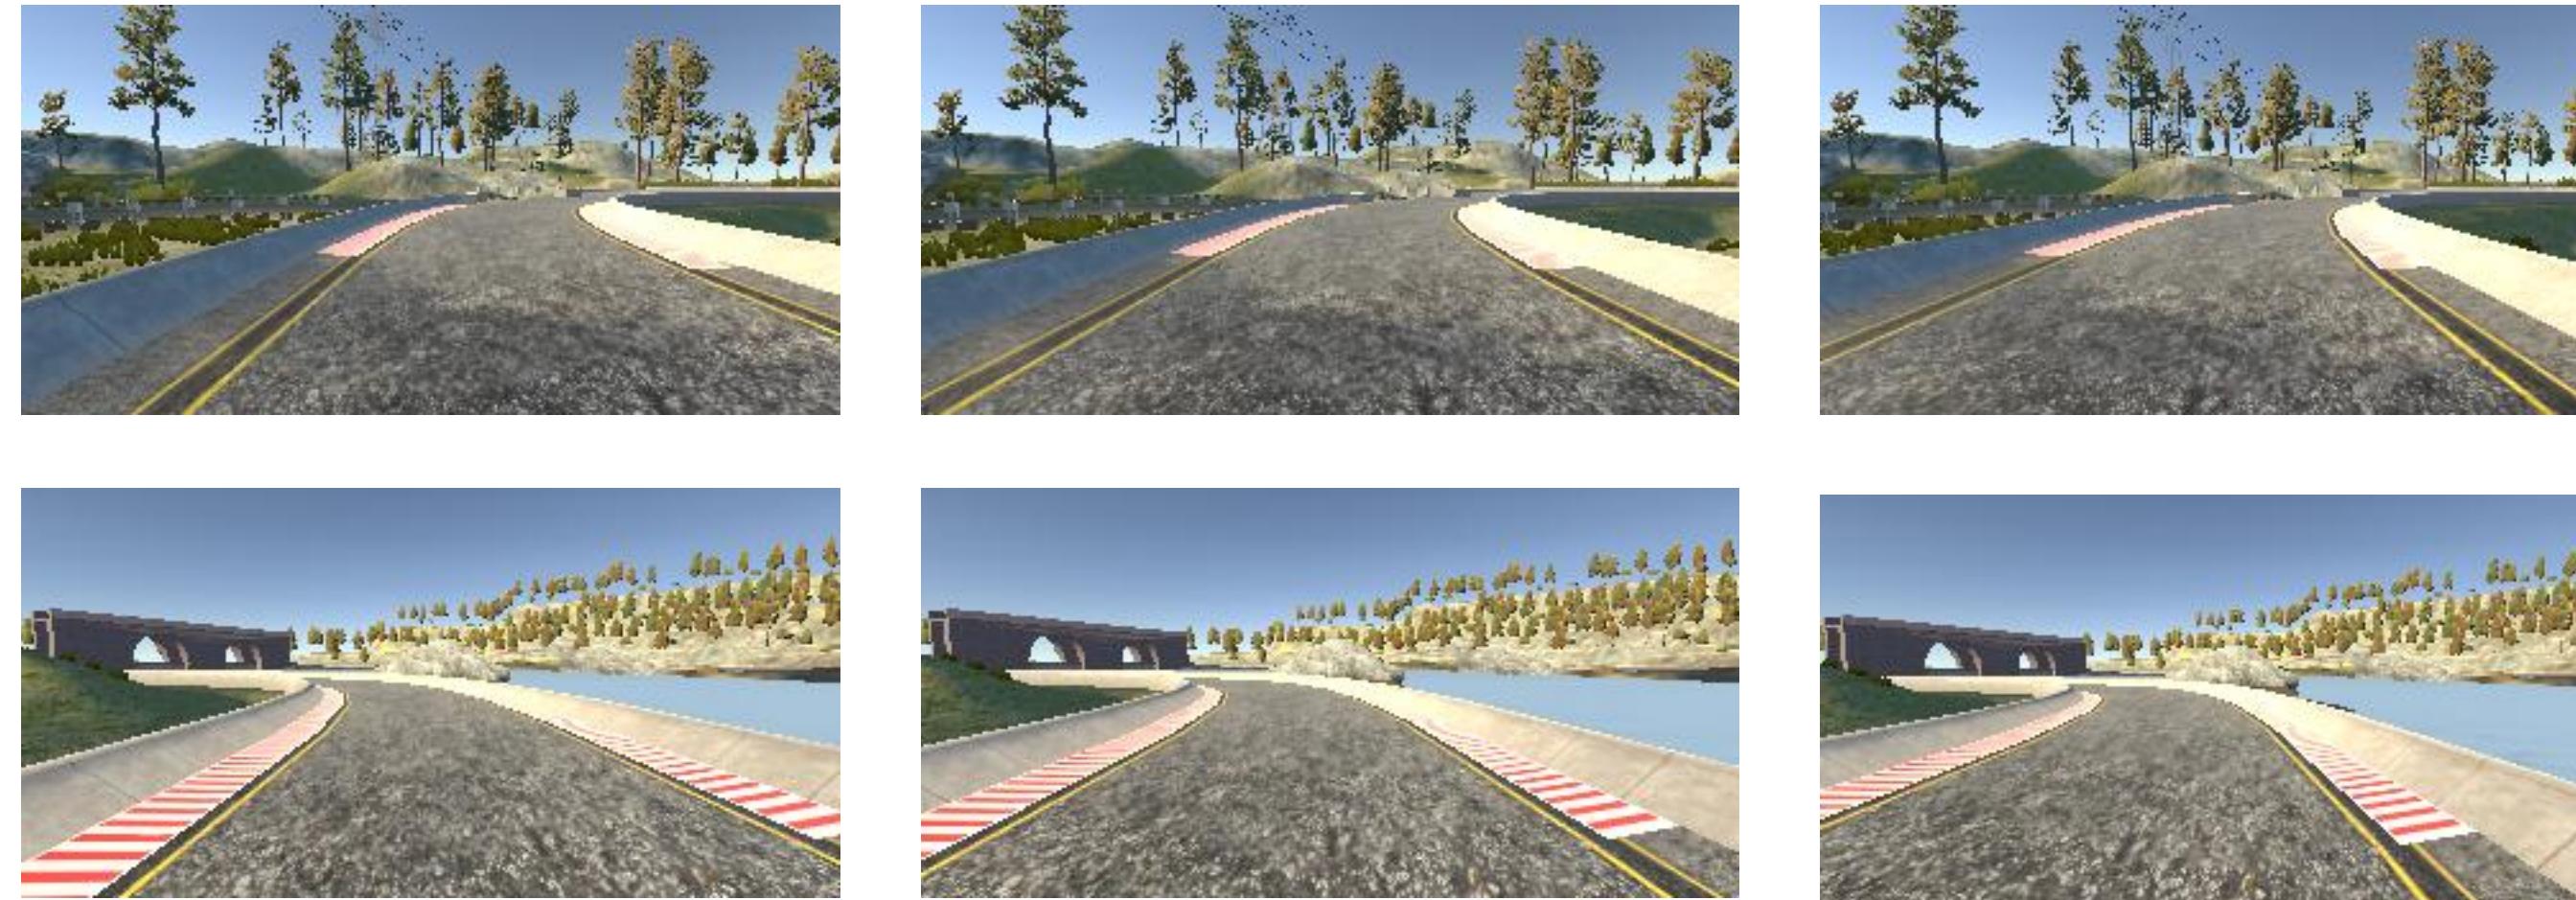
Self-Driving 3D car driven by a CNN

- The 3D environment in Unity



Dataset generation 1/2

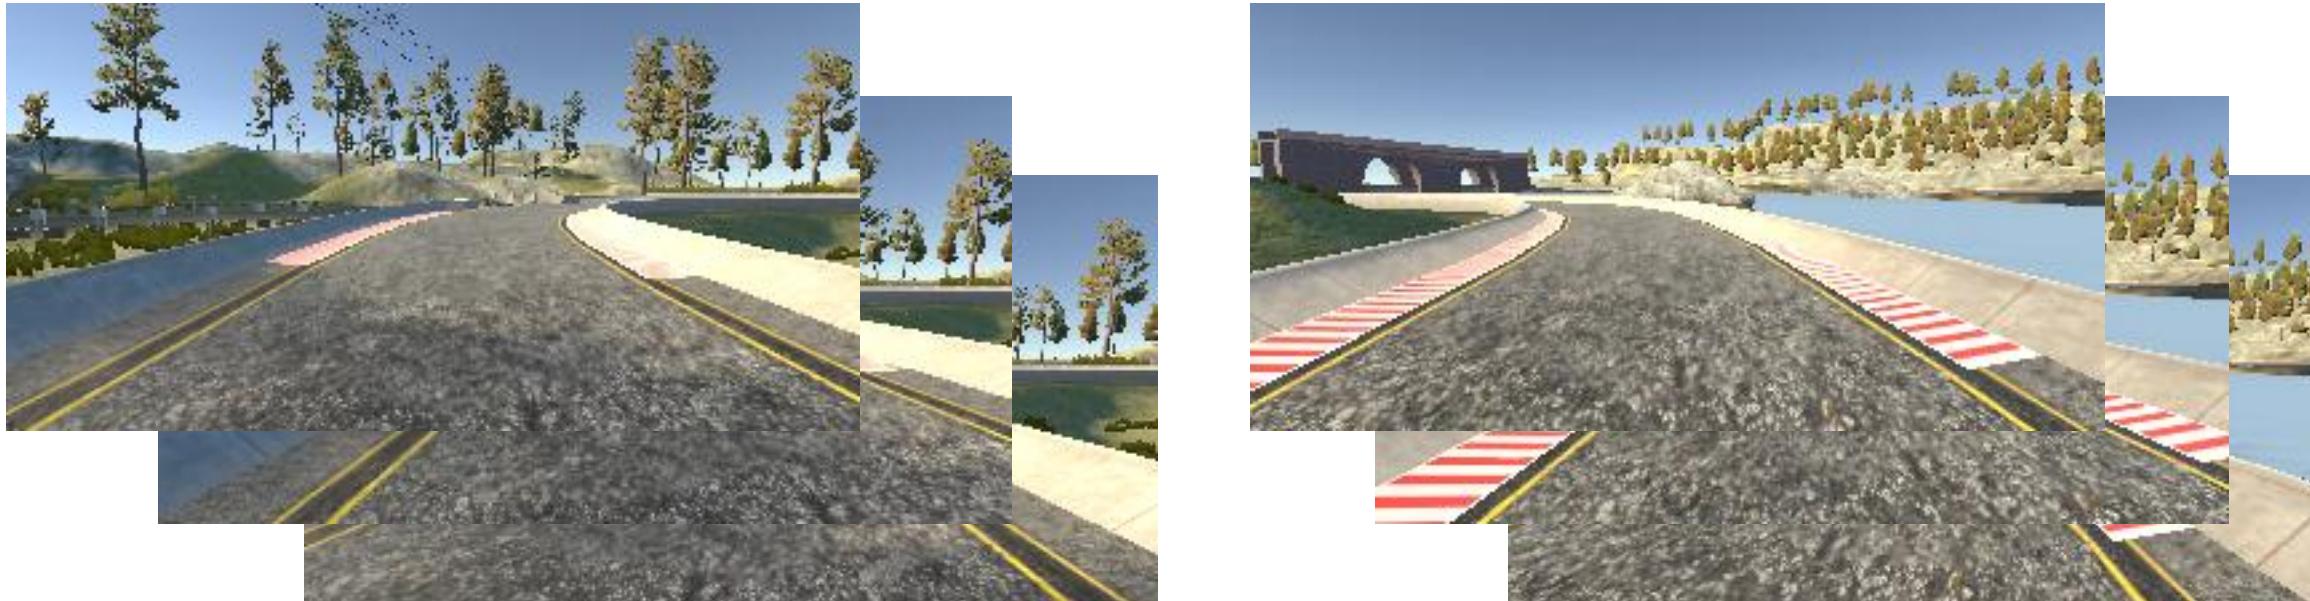
- Three **cameras** at constant distance on the x axis are bounded on the car
 - record the environment at 10 FPS
 - captured frames have size $320 \times 160 \times 3$



- We record the **steering angle** of the steering wheel
 - steering values in the non discrete range $[-1, 1]$ where $1 = \text{turn right}$, $0 = \text{go straight}$, $-1 = \text{turn left}$

Dataset generation 2/2

- The final dataset has the following components



4

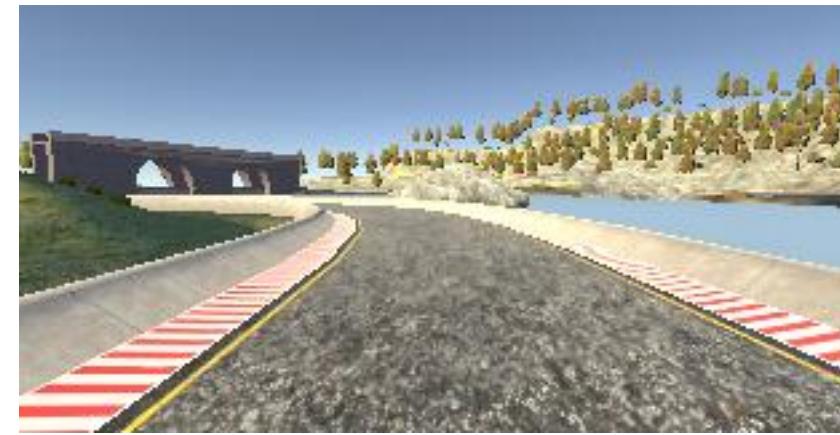


0, cam_mid_0.jpg, cam_lef_0.jpg, cam_rig_0.jpg, 0.7393
1, cam_mid_1.jpg, cam_lef_1.jpg, cam_rig_1.jpg, 0.7593
2, cam_mid_2.jpg, cam_lef_2.jpg, cam_rig_2.jpg, 0.7793
3, cam_mid_3.jpg, cam_lef_3.jpg, cam_rig_3.jpg, 0.7993
3, cam_mid_3.jpg, cam_lef_3.jpg, cam_rig_3.jpg, 0.7993

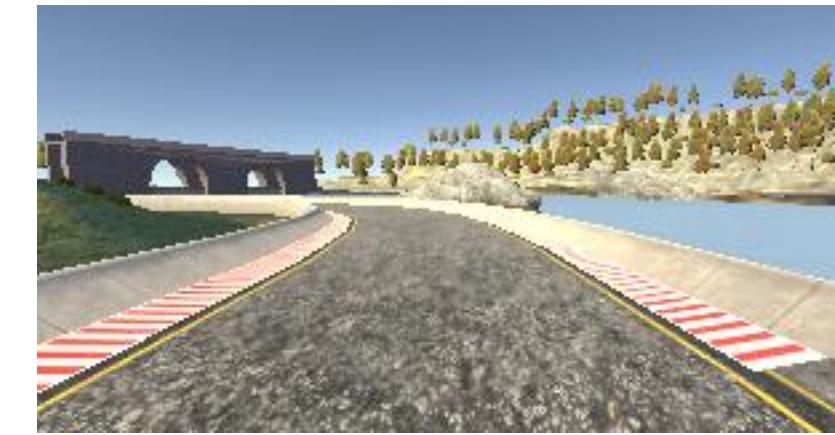
- Split dataset into **training set 80%** and **test set 20%**
 - Train the CNN with supervision by injecting it of **training images**, it produces a **steering angle**, prediction error is back propagated through the net

Dataset augmentation

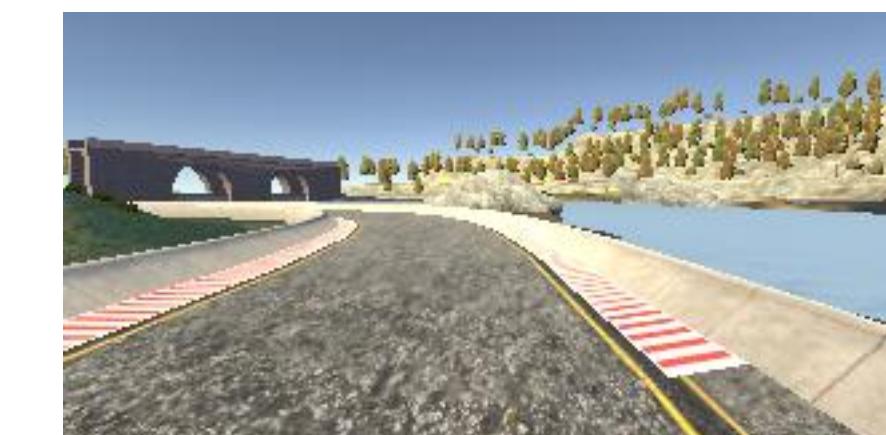
- Augmentation phase:
 - choose one only image between **left - mid - right** camera view + adjust steering angle



*former s.a.: -0.5
new s.a.: -0.2*



*former s.a.: -0.5
new s.a.: -0.5*



*former s.a.: -0.5
new s.a.: -0.8*

- randomly flip and slide images + adjust steering angle



former s.a.: +0.2



new s.a.: -0.2



former s.a.: +0.2



new s.a.: -0.1

Dataset preprocessing

- Preprocessing phase:
 - **crop** and **interpolate** the input image in order to remove the sky and useless street information



- apply **YUV filter** to make the image simpler for learning but same feeling of human eye



Keras implementation of the model

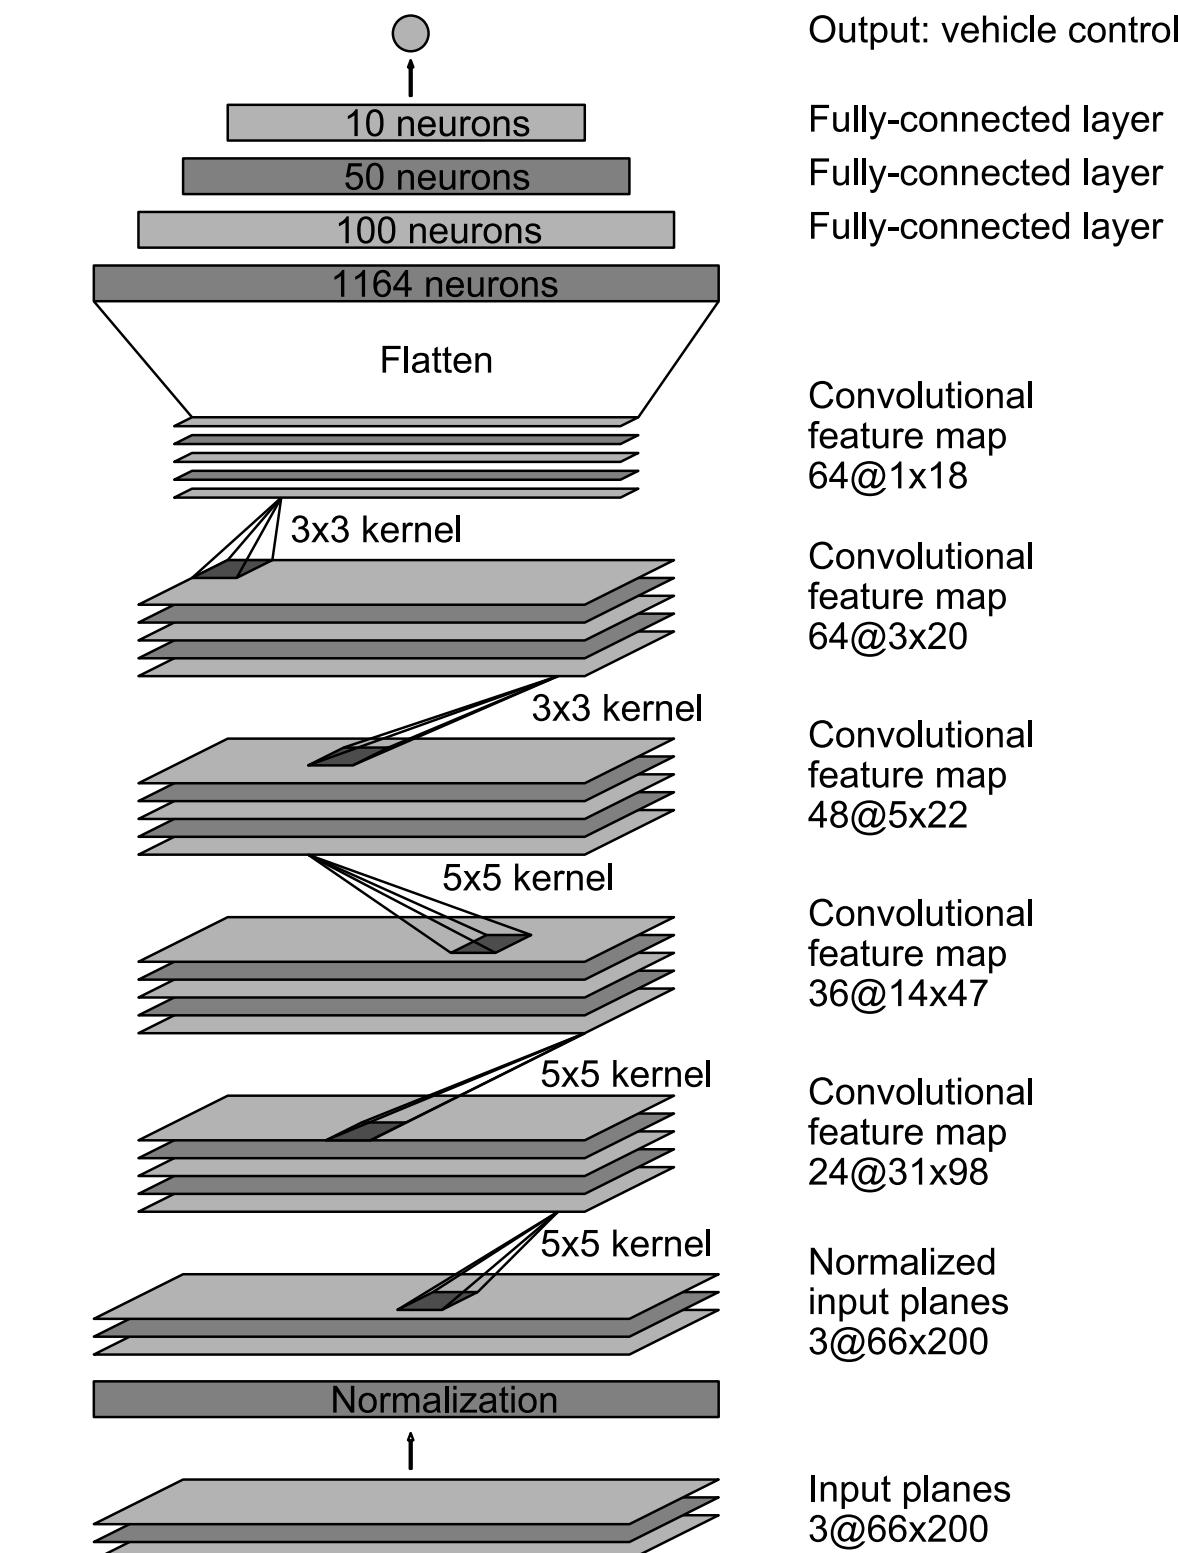
- We used **OpenCV** and **Keras** respectively for image preprocessing and neural implementation

```
def build_neural_model():
    """
    Build NVIDIA model:
    """

    INPUT_SHAPE = (IMAGE_HEIGHT, IMAGE_WIDTH, IMAGE_CHANNELS)

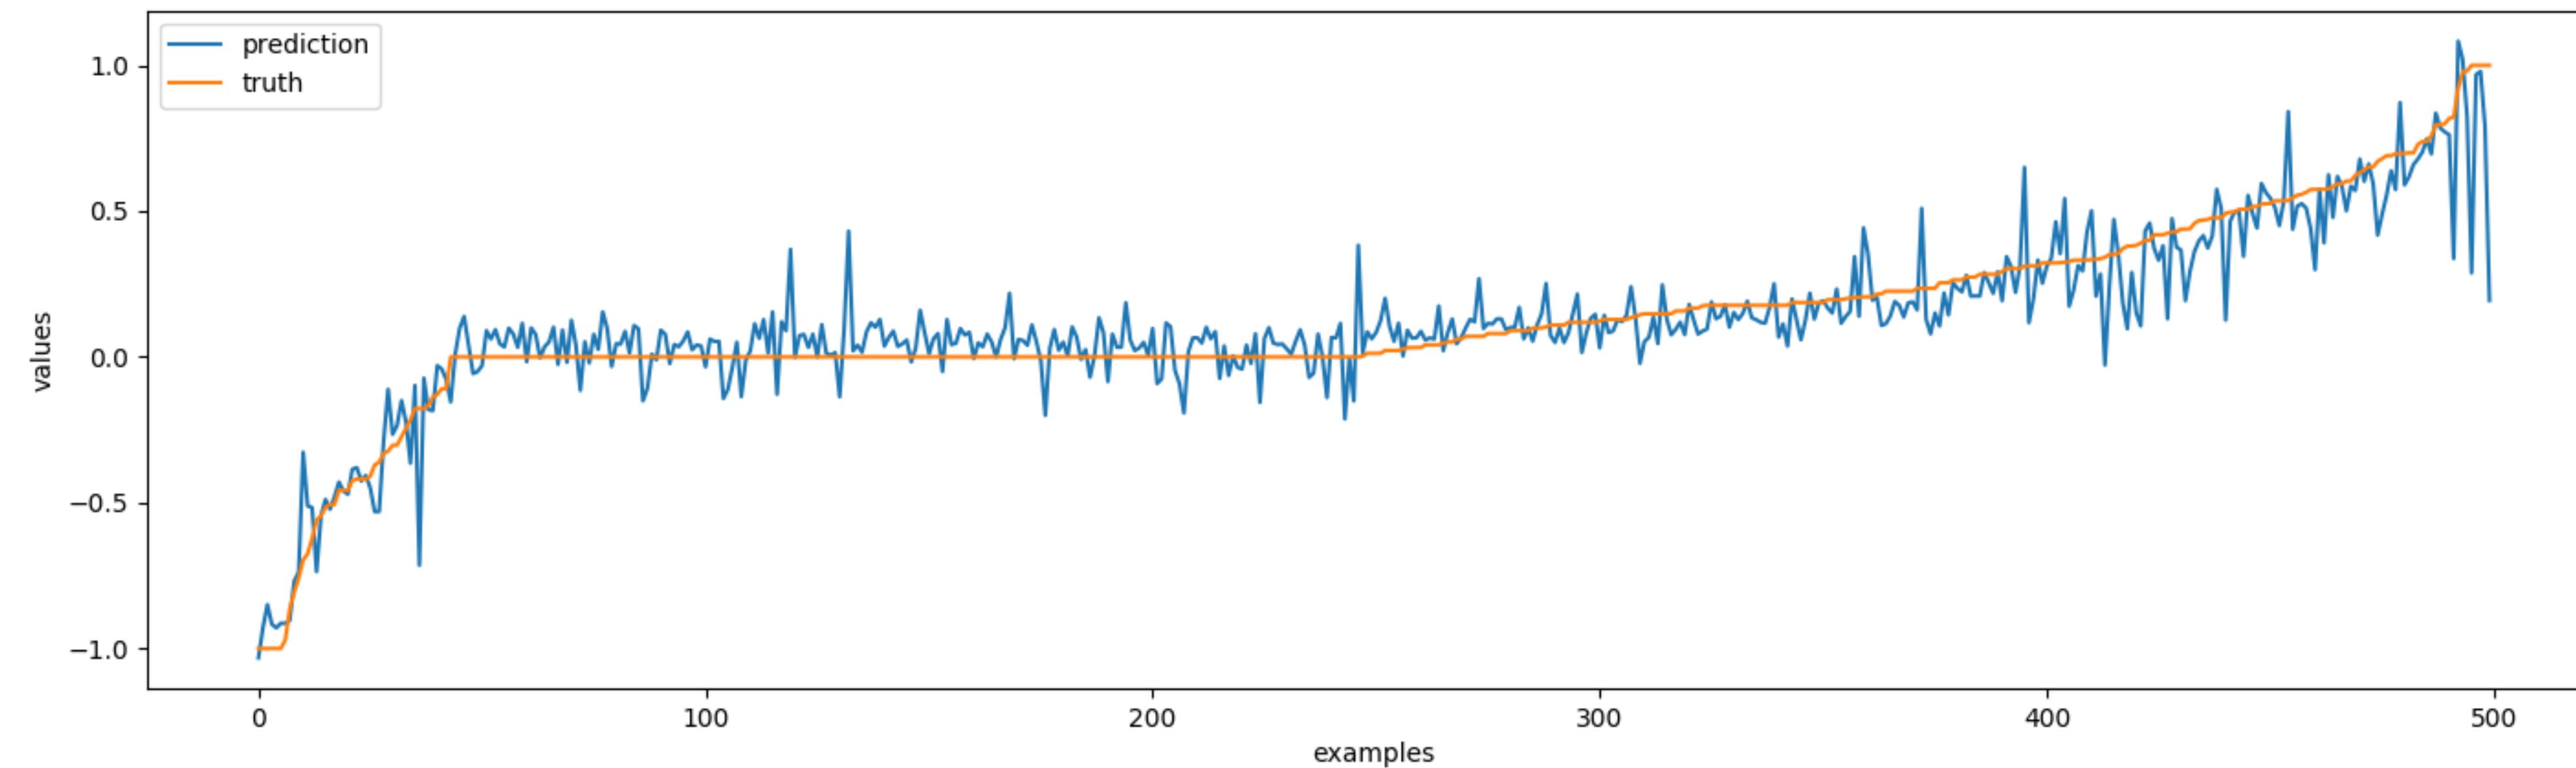
    model = Sequential()
    model.add(Lambda(lambda x: x / 127.5 - 1.0, input_shape=INPUT_SHAPE))
    model.add(Conv2D(24, 5, 5, activation='elu', subsample=(2, 2)))
    model.add(Conv2D(36, 5, 5, activation='elu', subsample=(2, 2)))
    model.add(Conv2D(48, 5, 5, activation='elu', subsample=(2, 2)))
    model.add(Conv2D(64, 3, 3, activation='elu'))
    model.add(Conv2D(64, 3, 3, activation='elu'))
    model.add(Dropout(0.5))
    model.add(Flatten())
    model.add(Dense(100, activation='elu'))
    model.add(Dense(50, activation='elu'))
    model.add(Dense(10, activation='elu'))
    model.add(Dense(1))
    model.summary()

    return model
```



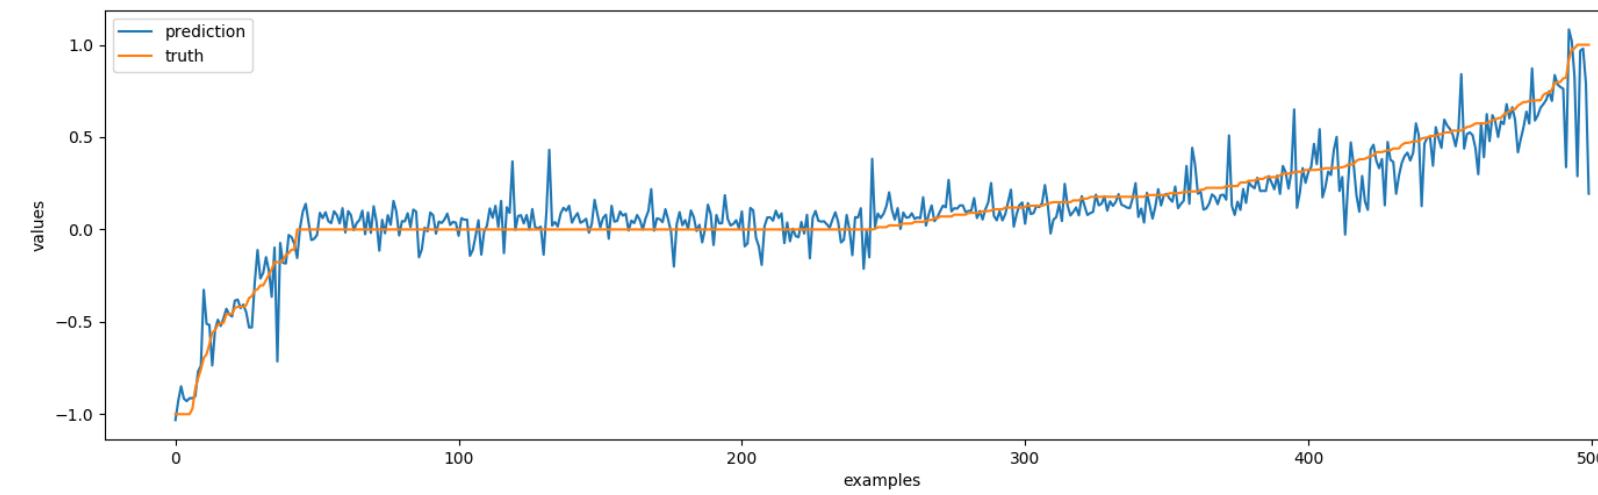
Evaluation

- Dataset generated by driving for **6 laps** using a **joypad**
- Training for **2 hours** (9000 instances, 300 epochs) with a GPU until loss reached its minimum
- Here's how well the CNN approximates the steering angle on 500 examples

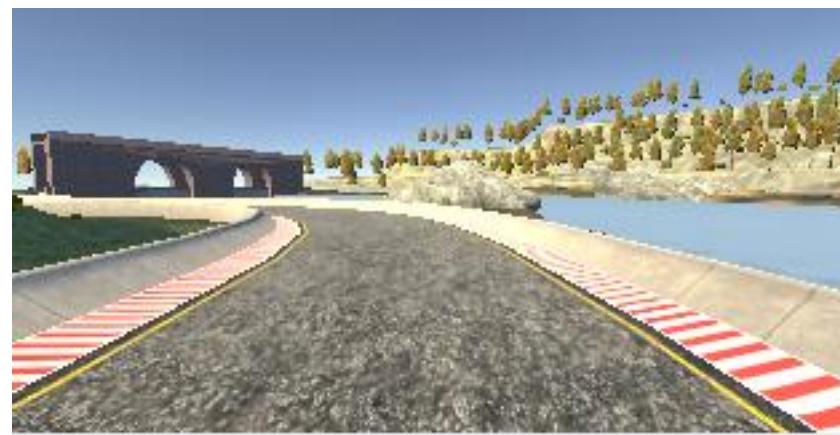


Evaluation

- Dataset generated by driving for **6 laps** using a **joypad**
- Training for **2 hours** (9000 instances, 300 epochs) with a GPU until loss reached its minimum
- Here's how well the CNN approximates the steering angle on 500 examples



- Here's what the network could see in an intermediate convolutional layers



Demo

Conclusion

- Make parameters tuning with a **grid search** approach
- Train on **diverse roads** with different light and weather conditions so make the model more resilient
- Integration with **sensors**, avoid **obstacles** and much complicated **decision making**
- **Clone now from GitHub here:** <https://github.com/matteoprata/self-deiving-car-cnn>



Thank you!